


Every node is trained at a balanced prior but deployed against a sub-1% node-local prior;
the two base learners miscalibrate in opposite directions

RFOED-RF: under-confident
(test-time generalization gap)



RFOED-NN: over-confident
(training-prior / threshold mismatch)

