


frequency order $\rightarrow$



Training augments link k with the *true* earlier labels: $[X | Y_{:, \ell_1}, \dots, Y_{:, \ell_{k-1}}]$



exposure bias

Inference must use the link's own *predicted* earlier labels: $[x | \hat{y}_{\ell_1}, \dots, \hat{y}_{\ell_{k-1}}]$

No example is ever removed; every link sees every document.